

CS 6364 Homework 6

October 26, 2021

Deadline for the first submission: **Nov-08-2021**.

All assignments **MUST** have your name, student ID, course name/number at the beginning of your documents.

Your homework **MUST** be submitted via Blackboard with file format and name convention as follows:

HW#_Name.writeup.pdf (for writing part)

HW#_Name.code.zip (for coding part)

If you have any questions, please contact me.

For the following questions, if you need a GPU to run, Google provide a free Jupyter notebook environment that requires no setup and runs entirely in the cloud. Here is the link:

https://colab.research.google.com/notebooks/welcome.ipynb?hl=enscrollTo=5fCEDCU_qrC0

A Markov decision Process is defined by $(\mathcal{S}, \mathcal{A}, \mathcal{R}, \mathbb{P}, \gamma)$, where \mathcal{S} denotes the set of possible states, \mathcal{A} denotes the set of possible actions, \mathcal{R} denotes distribution of reward given (state, action) pair, \mathbb{P} denotes transition probability, and γ is a discount factor. In this homework, given a simple 5×5 grid game (see below), the upper left (i.e. state 1-1) and lower right (i.e. state 5-5) corners are terminal states. Therefore, there are $5 \times 5 = 25$ states (i.e. $|\mathcal{S}| = 25$) and each is denoted as $s(i, j)$ where i and j represent i -th row and j -th column, respectively. Four possible actions are $\{right, left, up, down\}$. We set a negative reward $r = -5$ for each transition, a discount factor $\gamma = 0.7$, and the probability of an initial state $p(s_0)$ equals $\frac{1}{25}$. Our goal is to reach one of the terminal states (smiling states) in least number of actions. In other words, it aims to find the optimized policy π^* that maximized cumulative discounted reward. The initial Q table can be randomly defined by you.

- Q1 [Value iteration algorithm] Using the Bellman's equation, implement value iteration algorithm to iteratively update the Q table until convergence.
- Q2 [Deep Q-learning] This question functions the same as question 1. The only difference is to replace with a simple four-layer (i.e. three hidden layers and one output layer) fully-connected deep neural network to approximate the Q table. The three hidden layers contain 32, 8, and 16 neurons respectively and all applied ReLU activation functions. The output function is linear.
- Q3 [Deep Q-learning with experience replay, **bonus question**] This question functions the same as question 2, except that you need to implement the Deep Q-learning with experience replay. **Note that if you finish this bonus question, the points that you obtained in this HW will be doubled.**

Both two questions are programming assignments. Please use Pytorch for the implementation of the second one. Please submit your code along with initial and final Q -tables for both questions.

